

# Ri Hong

[riri.hong@gmail.com](mailto:riri.hong@gmail.com) | [github.com/Ri-Hong](https://github.com/Ri-Hong) | [linkedin.com/in/ririhong](https://linkedin.com/in/ririhong)

## SKILLS

**Languages:** Python, Java, JavaScript, C++, Go, TypeScript, HTML, CSS

**Technologies:** React, Next.js, Kubernetes, Docker, Terraform, gRPC, AWS, Linux, PostgreSQL, FastAPI

## EXPERIENCE

### Groq

Sept 2025 – Dec 2025

*Cloud Engineering Intern*

Toronto, ON

- Engineered multi-region infrastructure with **Kubernetes** and **Terraform** on GCP, enabling horizontally scalable systems and cutting deployment time by 24% for production workloads.
- Designed and launched an observability platform measuring engineering productivity for 400+ engineers, providing actionable insights through data-driven dashboards and analytics workflows.

### Base Power

Jan 2025 – Apr 2025

*Markets Infrastructure Engineer*

Austin, TX

- Led fullstack and infrastructure initiatives in a mission-critical environment, transitioning trading simulations to cloud-native execution using **Temporal** Cloud, enabling 1000s of auto-retrievable workflows and eliminating single-node failure risks.
- Discovered and resolved a performance bug in real-time market data transformation, reducing complexity from  $O(n^2)$  to  $O(n \log n)$ ; implemented non-blocking async publishing using **Go** routines, achieving a 32% speedup.
- Implemented **Protobuf** and **gRPC** for seamless communication between **Python** algorithm services and **Go** microservices, ensuring type safety and reducing serialization overhead by 20%.

### Trend Micro

Summer 2024

*Software Developer Intern*

Ottawa, ON

- Upgraded the legacy Deep Security Manager from JDK 8 to JDK 11, modernizing the **Java** codebase and enhancing compatibility with contemporary tools for over 250 million global customers.
- Revamped the **Jenkins** CI/CD pipeline to support JDK 11, achieving a 35% increase in automation efficiency and accelerating deployment timelines by 15%; refactored monolithic codebases into microservices, cutting deployment errors by 30%.

### Walnote.ai

Aug 2025 – Present

*Founder & CTO*

Toronto, ON

- Launched an AI platform combining GPT-5 with Manim to auto-generate explainer videos; processed 1,000+ animations with a **Celery** + **FastAPI** pipeline, accelerating generation by segmenting code and running distributed GPU rendering.

## PROJECTS

### DistilBERT Sentiment Analysis | PyTorch, GCP, Kubernetes, Terraform

Sept 2025

- Engineered a production-grade sentiment analysis service achieving 92.5% accuracy using DistilBERT, deployed on GKE with **Terraform**; optimized training with mixed precision and distributed GPU training, reducing training time by 40%.

### Neural Style Transfer Engine | PyTorch, CUDA, FastAPI, Docker

Sept 2025

- Engineered a high-performance Neural Style Transfer system with custom **CUDA** kernels, achieving 8x faster style transfer; built a scalable REST API with **FastAPI** for production deployment.

### Trasee | Pyodide, React, TypeScript, Vite

Oct 2025

- Built a real-time **Python** code visualizer that intelligently recognizes data structures and renders them interactively in the browser; engineered a two-phase static + runtime analysis pipeline using **Python**'s ast module and sys.settrace within **Pyodide** WebAssembly.

## EDUCATION

### University of Waterloo

2022 – 2027

*Bachelor of Computer Science (Co-op) · GPA: 3.9/4.0*

- Relevant coursework: Algorithms, Data Structures, Object-Oriented Programming, Databases, AI, ML